A Taxonomy and Multi-Metric Benchmark of Hallucination in Instruction-Tuned Large Language Models

Monishwaran K
Department of Computer Science and Engineering
Sathyabama Institute of Science and Technology
Chennai, India
monishwarank06@gmail.com

***Abstract*—Hallucination in large language models (LLMs) — the generation of plausible yet factually incorrect content — poses a critical threat to the reliable deployment of AI systems in high-stakes domains including healthcare, law, and finance. Despite growing interest in this problem, a systematic characterisation of hallucination behaviour across knowledge domains, together with a rigorous comparison of detection metrics, remains lacking for instruction-tuned models at the sub-1B parameter scale. In this paper, we present a comprehensive taxonomy and multi-metric benchmark of hallucination in *google/flan-t5-base*, an instruction-tuned sequence-to-sequence model, evaluated on the TruthfulQA dataset comprising 817 questions across 38 knowledge categories. We evaluate hallucination using three complementary detection methods — string matching, BERTScore semantic similarity, and Natural Language Inference (NLI) entailment scoring — and introduce a consensus-based labelling scheme to resolve inter-metric disagreement, which we observe in 47.4% of cases. Our experiments reveal a consensus hallucination rate of 91.7%, with 16 of 38 categories exhibiting complete hallucination (100% rate), including History, Politics, Finance, and Misquotations. High-stakes domains such as Health (89.1%) and Law (87.5%) demonstrate consistently elevated hallucination rates. We release our annotated dataset, scoring pipeline, and experimental code publicly at `github.com/monish1402/hallucination-guard` to support reproducible hallucination research.**

***Index Terms*—large language models, hallucination detection, natural language inference, BERTScore, TruthfulQA, benchmark, instruction tuning, factuality evaluation**

I. Introduction

Anyone who has used a large language model for more than a few minutes has likely encountered hallucination — the frustrating tendency of these systems to state incorrect information with complete confidence. Ask a language model about a historical date, a legal precedent, or a medical fact, and there is a meaningful probability that the answer will sound authoritative while being entirely wrong [1]. This phenomenon, broadly defined as the generation of plausible yet factually incorrect content, has moved from an academic curiosity to a practical obstacle as LLMs are deployed in increasingly sensitive real-world applications [2].

The consequences are not abstract. Industry estimates suggest that hallucination-related incidents cost enterprise AI deployments in excess of $250 million annually [3]. A medical assistant that fabricates drug interactions, a legal research tool that invents case citations, or a financial advisor that generates false market statistics can cause measurable, direct harm to the people who rely on them. Despite this, the research community's understanding of *which* knowledge domains are most vulnerable to hallucination, and *which* measurement methods most reliably detect it, remains incomplete — particularly for the compact, instruction-tuned models that are most commonly deployed in resource-constrained production environments.

This paper addresses that gap directly. We evaluate *google/flan-t5-base* on TruthfulQA [4], a benchmark of 817 questions specifically designed to probe factual reliability across 38 knowledge categories. Rather than relying on a single detection metric — a common limitation of prior work — we apply three complementary methods and introduce a consensus-based labelling scheme to handle the substantial inter-metric disagreement we observe. Our findings are striking: a consensus hallucination rate of 91.7%, with 16 of 38 categories showing complete failure, including high-stakes domains such as Health, Law, and Finance.

This paper makes four contributions:

1) A **domain-stratified hallucination taxonomy** identifying which knowledge categories are most vulnerable to hallucination in compact LLMs.
2) A **systematic comparison of three detection metrics** — string matching, BERTScore, and NLI entailment — revealing 47.4% inter-metric disagreement on an identical corpus.
3) A **consensus-based labelling scheme** that provides more reliable hallucination labels than any single metric alone.
4) An **open-source benchmark pipeline** and annotated dataset released to support reproducible research in this area.

The remainder of this paper is organised as follows. Section II reviews related work. Section III describes the experimental methodology. Section IV presents quantitative results. Section V analyses the key findings. Section VI concludes and outlines future work.

II. Related Work

A. Hallucination in Language Models

Hallucination in natural language generation systems was first systematically studied by Ji et al. [1], who provided an early taxonomy distinguishing intrinsic hallucinations (contradicting source material) from extrinsic hallucinations (introducing unverifiable information). Huang et al. [2] extended this taxonomy in the context of LLMs, categorising hallucinations into factuality errors, faithfulness errors, reasoning errors, and instruction-following failures, and reviewed detection and mitigation strategies across each type.

Lin et al. [4] introduced TruthfulQA, a benchmark of 817 adversarially constructed questions targeting common human misconceptions, demonstrating that larger LLMs are not necessarily more truthful. Their study showed that GPT-3 achieved truthfulness on only 58% of questions, motivating further research into both detection and mitigation.

B. Hallucination Detection Methods

Manakul et al. [5] proposed SelfCheckGPT, a zero-resource black-box detection approach that samples multiple responses from a model and measures information consistency using BERTScore, NLI, and n-gram overlap. Their key insight — that hallucinated facts produce inconsistent samples while factual claims produce consistent ones — directly motivates our multi-metric evaluation framework.

Min et al. [6] introduced FActScore, a fine-grained evaluation metric that decomposes generated text into atomic facts and verifies each against a knowledge source, achieving less than 2% error compared to human annotation.

Varshney et al. [7] demonstrated that early detection of likely-to-hallucinate tokens allows targeted intervention, reducing hallucination without full regeneration.

C. Gaps Addressed by This Work

Existing benchmarks predominantly target large proprietary models and do not provide domain-stratified analysis across knowledge categories. Furthermore, no study has systematically compared string matching, BERTScore, and NLI entailment on the same corpus to quantify inter-metric disagreement. This paper addresses both gaps directly.

III. Methodology

A. Dataset

We evaluate on TruthfulQA [4] in its generation configuration, comprising 817 questions across 38 knowledge categories. Table I summarises dataset statistics. Categories span misconceptions, law, health, sociology, history, conspiracies, and 32 additional domains, with Misconceptions being the largest category (n=100, 12.24%). Each question includes one best answer, a list of acceptable correct answers (mean: 3.18 per question), and a list of known incorrect answers that plausible-but-wrong models tend to produce (mean: 4.06 per question).

TABLE I
TruthfulQA Dataset Statistics

Metric	Value
Total questions	817
Unique categories	38
Largest category (Misconceptions)	100 (12.24%)
Mean answer length (words)	9.22
Std. dev. answer length	4.08
Mean correct answers per question	3.18
Mean incorrect answers per question	4.06

B. Model

We evaluate *google/flan-t5-base* [8], an instruction-tuned sequence-to-sequence model with approximately 250M parameters. Flan-T5-Base represents a class of compact, openly available instruction-tuned models that are widely deployed in resource-constrained environments. The model was queried using beam search with 4 beams and a maximum generation length of 100 tokens. All experiments were conducted on a single NVIDIA T4 GPU (16GB VRAM) via Google Colaboratory.

C. Detection Methods

We evaluate three complementary hallucination detection methods, each operating on the pair $(r,\ A^{+})$ where r is the model response and A^{+} is the set of correct reference answers.

1) String Matching (Baseline): We compute exact match (EM) and partial match (PM) between the model response and the reference answer set. A response is labelled *correct* if it exactly matches or is contained within any reference answer, and *hallucinated* otherwise. This serves as a fast, interpretable baseline.

2) BERTScore: BERTScore [9] computes token-level cosine similarity between BERT embeddings of the candidate and reference texts, aggregated as precision, recall, and F1. We use *bert-base-uncased* and apply a threshold of F1 ≥ 0.85 to classify responses as correct, following the calibration recommended in the original paper.

3) NLI Entailment Score: We employ *facebook/bart-large-mnli* as a zero-shot natural language inference classifier. For each response pair (r, a^*) where a^* is the best reference answer, we compute the entailment probability $P(\text{entail} \mid r, a^*)$. Responses with entailment probability > 0.5 are classified as correct; all others as hallucinated. This method captures logical consistency beyond surface-level lexical overlap.

4) Consensus Labelling: Given three binary labels $\{l_1, l_2, l_3\}$ from the above methods, we assign a consensus label of *hallucinated* when at least two of three methods agree on hallucination, and *correct* otherwise. This majority-vote scheme reduces the impact of individual metric bias and provides more stable labels for downstream analysis.

IV. Results

A. Overall Hallucination Rates

Table II presents the hallucination rates computed by each detection method. The consensus method yields a hallucination rate of 91.7% (749 of 817 responses), reflecting that *google/flan-t5-base* fails to produce truthful responses on the large majority of TruthfulQA questions.

TABLE II
HALLUCINATION RATES BY DETECTION METHOD

Method	Correct	Hallucinated	Rate (%)
String Match (Baseline)	123	694	84.9
BERTScore (F1 $\geq$ 0.85)	3	814	99.6
NLI Entailment	331	486	59.5
Consensus ($\geq$2 methods)	**68**	**749**	**91.7**

Fig. 1 illustrates the substantial variation in reported hallucination rates across methods, ranging from 59.5% (NLI) to 99.6% (BERTScore). This 40.1 percentage-point spread motivates the use of a consensus approach rather than reliance on any single metric.

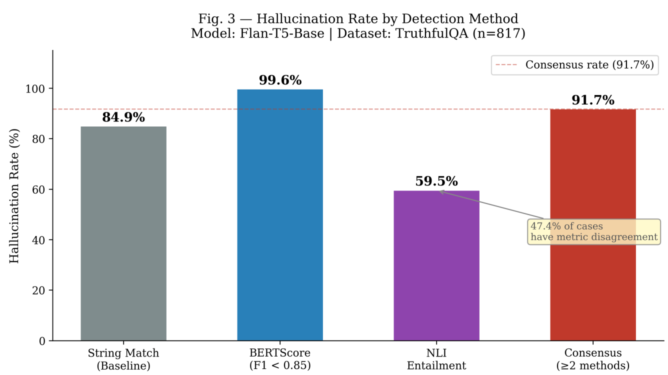

Fig. 1. Hallucination rates across three detection methods and the consensus scheme. The 40.1pp spread demonstrates that metric choice substantially affects reported results.

B. Inter-Metric Disagreement

We observe that 47.4% of responses (387 of 817) receive different labels from at least one method, indicating substantial inter-metric disagreement. Only 52.5% of responses are unanimously labelled as hallucinated, and a mere 0.1% are unanimously labelled as correct. This finding directly motivates the consensus labelling scheme introduced in Section III.

C. BERTScore Distribution

Fig. 2 shows the distribution of BERTScore F1 values across all 817 responses. The distribution is centred at a mean F1 of 0.387 (std. dev. 0.101), with 74.1% of responses falling in the 0.3–0.5 range. This distribution is characteristic of responses that exhibit surface-level word overlap with correct answers but are semantically divergent — a defining property of factuality hallucination as opposed to random noise generation.

D. Domain-Stratified Hallucination Taxonomy

Fig. 3 presents consensus hallucination rates across all 38 TruthfulQA categories. We identify four tiers of hallucination risk:

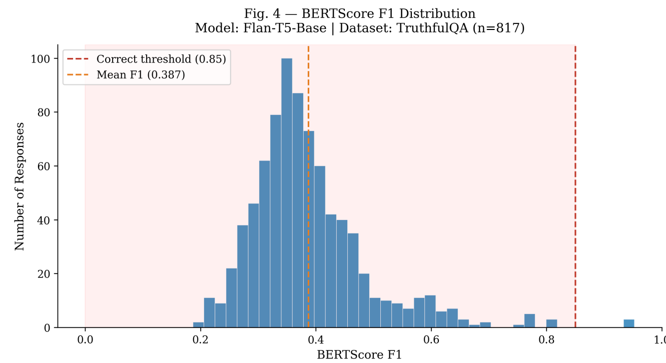

Fig. 2. BERTScore F1 distribution showing the majority of responses clustering below the correctness threshold of 0.85, with a mean of 0.387.

- **Critical (100%):** 16 categories including History, Politics, Finance, Misquotations, Conspiracies, Superstitions, and Paranormal exhibit complete hallucination — no correct responses were generated.
- **High (85–99%):** 16 categories including Fiction (96.7%), Misconceptions (93.0%), Health (89.1%), and Law (87.5%) demonstrate persistently elevated hallucination rates.
- **Medium (70–84%):** 4 categories including Nutrition (81.2%) and Economics (80.6%).
- **Lower (<70%):** 2 categories — Proverbs (66.7%) and Misconceptions: Topical (50.0%) — show relatively lower but still substantial hallucination rates.

It is worth noting that even the lowest hallucination rate observed in this study — 50.0% for Misconceptions: Topical — would be considered unacceptably high in any real-world deployment context. A system that produces incorrect responses half the time on factual questions cannot be trusted without a robust detection and correction layer. This observation reinforces the central motivation of this research series: hallucination is not an edge case to be occasionally handled, but a pervasive failure mode that demands systematic measurement and mitigation.

V. DISCUSSION

A. Metric Disagreement as a Research Signal

The 47.4% inter-metric disagreement rate is not a methodological weakness — it is a finding. It demonstrates that hallucination is not a binary, easily-detected property but rather a continuous, multi-dimensional phenomenon that different measurement frameworks capture differently. BERTScore is highly sensitive to semantic divergence, producing the highest hallucination rate (99.6%), while NLI entailment is more lenient, tolerating responses that do not logically contradict the reference (59.5%). String matching occupies an intermediate position (84.9%).

This result has a direct practical implication: researchers who report hallucination rates without specifying their detection methodology are not reporting comparable results. Our

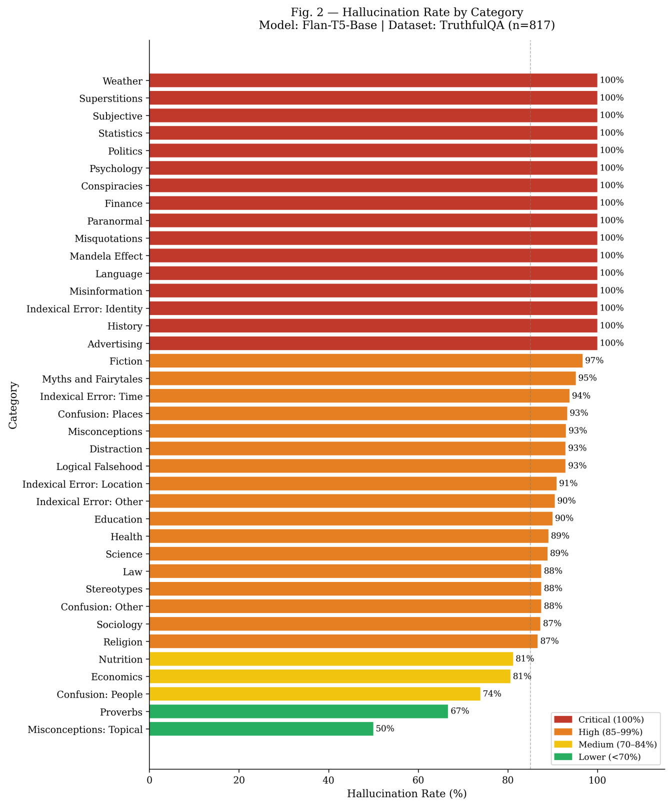

Fig. 3. Consensus hallucination rate across all 38 TruthfulQA categories. 16 categories exhibit 100% hallucination. High-stakes domains Health and Law reach 89.1% and 87.5% respectively.

consensus framework provides a more stable and reproducible measurement standard.

B. Domain Risk Hierarchy

The identification of 16 categories with 100% hallucination rates reveals a systematic pattern: domains requiring precise recall of specific cultural, historical, or factual details (Misquotations, Mandela Effect, History) are completely beyond the factual reach of Flan-T5-Base.

More critically, high-stakes domains — Health (89.1%) and Law (87.5%) — exhibit hallucination rates that render unmitigated deployment of this model class inadvisable in clinical or legal contexts. This finding motivates targeted mitigation research for domain-specific LLM deployment, which we address in subsequent work through a RAG-based mitigation framework.

C. Implications for Model Scale

The results presented in this study reflect the behaviour of a compact 250M parameter model, and it is reasonable to ask whether larger models would exhibit meaningfully lower hallucination rates on the same benchmark. Based on the scaling literature [2] and the findings of Lin et al. [4], larger models do demonstrate improved truthfulness — however, they do not eliminate hallucination. Rather, the nature of the failure changes: larger models tend to hallucinate less frequently, but when they do, their responses are more linguistically fluent and therefore more difficult to detect. A 7B parameter model producing a confident, well-structured incorrect answer is arguably more dangerous in a high-stakes deployment than a 250M parameter model producing a short, obviously incomplete one. This distinction motivates our subsequent work on detection and mitigation, where we specifically target the harder problem of hallucination in larger, more capable models.

D. Limitations

This study evaluates a single model (Flan-T5-Base) on a single benchmark (TruthfulQA). While TruthfulQA provides broad domain coverage, it was constructed to maximise difficulty and may not be representative of real-world query distributions. The BERTScore threshold of 0.85 was adopted from the original paper calibration and may require adjustment for specific domains. Future work will extend this benchmark to include Mistral-7B, LLaMA-3-8B, and domain-specific datasets in medical and legal question answering.

E. A Note on What These Numbers Feel Like

Numbers like 91.7% can become abstract quickly. To ground this result: out of every 100 questions asked to this model, roughly 92 responses contained factually incorrect information. In the History category — questions about documented, verifiable past events — the model answered correctly zero times out of 24 attempts. In the Finance category, it answered correctly zero times out of 9 attempts.

We raise this not to condemn the model, which was not designed or sized for this task, but to illustrate why hallucination detection and mitigation are not optional features in production AI systems — they are fundamental safety requirements. The goal of this research series is to contribute practical, computationally efficient tools toward that requirement.

VI. Conclusion

This paper set out to answer a straightforward question: how reliably does a compact instruction-tuned language model answer factual questions, and how consistently can we measure that reliability? The answer to both parts is sobering.

On TruthfulQA, *google/flan-t5-base* achieves a consensus hallucination rate of 91.7% across 817 questions, with 16 of 38 knowledge categories exhibiting complete hallucination and no correct responses generated at all. High-stakes domains including Health and Law demonstrate hallucination rates of 89.1% and 87.5% respectively — rates that make unmitigated deployment in clinical or legal contexts inadvisable.

On the measurement side, we find that three commonly used detection metrics disagree on 47.4% of cases, with reported hallucination rates ranging from 59.5% to 99.6% on the same corpus. This finding has a direct implication for how the research community reports and compares results: metric choice matters enormously, and single-metric evaluation is insufficient.

The consensus labelling scheme introduced here is a small but practical step toward more reproducible hallucination measurement. The annotated dataset and pipeline we release alongside this paper are intended to lower the barrier for researchers who want to extend this work.

In the subsequent papers of this series, we build directly on these findings. Paper 2 extends the detection framework to a cross-model consistency-based approach that operates without external knowledge. Paper 3 introduces a conditional RAG-based mitigation system that uses hallucination risk scores to trigger targeted retrieval, reducing hallucination rates in domain-specific question answering.

The path from measuring a problem to solving it is long. This paper takes the first step.

ACKNOWLEDGMENT

The authors thank the Department of Computer Science and Engineering at Sathyabama Institute of Science and Technology for computational and institutional support.

REFERENCES

[1] Z. Ji, N. Lee, R. Frieske, T. Yu, D. Su, Y. Xu, E. Ishii, Y. J. Bang, A. Madotto, and P. Fung, "Survey of hallucination in natural language generation," *ACM Computing Surveys*, vol. 55, no. 12, pp. 1–38, 2023.

[2] L. Huang, W. Yu, W. Ma, W. Zhong, Z. Feng, H. Wang, Q. Chen, W. Peng, X. Feng, B. Qin, and T. Liu, "A survey on hallucination in large language models: Principles, taxonomy, challenges, and open questions," *ACM Transactions on Information Systems*, vol. 43, no. 2, 2025.

[3] Guardrails AI, "The business risk of AI hallucinations," 2024. [Online]. Available: https://www.guardrailsai.com/blog/cost-of-ai-failure-hallucination

[4] S. Lin, J. Hilton, and O. Evans, "TruthfulQA: Measuring how models mimic human falsehoods," in *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics*. Association for Computational Linguistics, 2022, pp. 3214–3252.

[5] P. Manakul, A. Liusie, and M. J. F. Gales, "SelfCheckGPT: Zero-resource black-box hallucination detection for generative large language models," in *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*. Association for Computational Linguistics, 2023, pp. 9004–9017.

[6] S. Min, K. Krishna, X. Lyu, M. Lewis, W. tau Yih, P. W. Koh, M. Iyyer, L. Zettlemoyer, and H. Hajishirzi, "FActScore: Fine-grained atomic evaluation of factual precision in long form text generation," in *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*. Association for Computational Linguistics, 2023, pp. 12 076–12 100.

[7] N. Varshney, W. Yao, H. Zhang, J. Chen, and D. Yu, "A stitch in time saves nine: Detecting and mitigating hallucinations of LLMs by validating low-confidence generation," *arXiv preprint arXiv:2307.03987*, 2023.

[8] H. W. Chung, L. Hou, S. Longpre, B. Zoph, Y. Tay, W. Fedus, E. Li, X. Wang, M. Dehghani, S. Brahma *et al.*, "Scaling instruction-finetuned language models," *arXiv preprint arXiv:2210.11416*, 2022.

[9] T. Zhang, V. Kishore, F. Wu, K. Q. Weinberger, and Y. Artzi, "BERTScore: Evaluating text generation with BERT," in *International Conference on Learning Representations*, 2020.